

NAME : EDE RENUKA MADHAV

REG NO: 21BCT0244

VIT VELLORE

SLOT: 6:00 PM TO 8:00 PM

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
```

Load the Dataset

```
df=pd.read_csv("/content/winequality-red.csv")
df.head()
```

	fixed acidity	volatile acidity	citric acid	residual sugar
0	7.4	0.70	0.00	1.9
1	7.8	0.88	0.00	2.6
2	7.8	0.76	0.04	2.3
3	11.2	0.28	0.56	1.9
4	7.4	0.70	0.00	1.9

	free sulfur dioxide	total sulfur dioxide	density	pH	sulphates
0	11.0	34.0	0.9978	3.51	0.56
1	25.0	67.0	0.9968	3.20	0.68
2	15.0	54.0	0.9970	3.26	0.65
3	17.0	60.0	0.9980	3.16	0.58
4	11.0	34.0	0.9978	3.51	0.56

	alcohol	quality
0	9.4	5
1	9.8	5
2	9.8	5
3	9.8	6
4	9.4	5

```
df.shape
```

```
(1599, 12)
```

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 1599 entries, 0 to 1598
```

```
Data columns (total 12 columns):
```

#	Column	Non-Null Count	Dtype
0	fixed acidity	1599 non-null	float64
1	volatile acidity	1599 non-null	float64
2	citric acid	1599 non-null	float64
3	residual sugar	1599 non-null	float64
4	chlorides	1599 non-null	float64
5	free sulfur dioxide	1599 non-null	float64
6	total sulfur dioxide	1599 non-null	float64
7	density	1599 non-null	float64
8	pH	1599 non-null	float64
9	sulphates	1599 non-null	float64
10	alcohol	1599 non-null	float64
11	quality	1599 non-null	int64

```
dtypes: float64(11), int64(1)
```

```
memory usage: 150.0 KB
```

Data preprocessing

VISUALIZATIONS

1. UNIVARIATE ANALYSIS

```
sns.distplot(df.alcohol)
```

```
<ipython-input-9-cc0e16fd78a8>:1: UserWarning:
```

```
`distplot` is a deprecated function and will be removed in seaborn
```


v0.14.0.

Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).

For a guide to updating your code to use the new functions, please see <https://gist.github.com/mwaskom/de44147ed2974457ad6372750bbe5751>

```
sns.distplot(df.alcohol)
```


```
<Axes: xlabel='alcohol', ylabel='Density'>
```



2. BIVARIATE ANALYSIS

```
sns.lineplot(x = df.alcohol ,y = df.sulphates)
```


```
<Axes: xlabel='alcohol', ylabel='sulphates'>
```



3.MULTIVARIATE ANALYSIS

```
sns.pairplot(df)
```

```
<seaborn.axisgrid.PairGrid at 0x7f5c1ddd95d0>
```



STATISTICAL DATA

```
df.describe()
```

	fixed acidity	volatile acidity	citric acid	residual sugar	\
count	1599.000000	1599.000000	1599.000000	1599.000000	
mean	8.319637	0.527821	0.270976	2.538806	
std	1.741096	0.179060	0.194801	1.409928	
min	4.600000	0.120000	0.000000	0.900000	
25%	7.100000	0.390000	0.090000	1.900000	
50%	7.900000	0.520000	0.260000	2.200000	

75%	9.200000	0.640000	0.420000	2.600000
max	15.900000	1.580000	1.000000	15.500000
	chlorides	free sulfur dioxide	total sulfur dioxide	
density \				
count	1599.000000	1599.000000	1599.000000	
1599.000000				
mean	0.087467	15.874922	46.467792	
0.996747				
std	0.047065	10.460157	32.895324	
0.001887				
min	0.012000	1.000000	6.000000	
0.990070				
25%	0.070000	7.000000	22.000000	
0.995600				
50%	0.079000	14.000000	38.000000	
0.996750				
75%	0.090000	21.000000	62.000000	
0.997835				
max	0.611000	72.000000	289.000000	
1.003690				
	pH	sulphates	alcohol	quality
count	1599.000000	1599.000000	1599.000000	1599.000000
mean	3.311113	0.658149	10.422983	5.636023
std	0.154386	0.169507	1.065668	0.807569
min	2.740000	0.330000	8.400000	3.000000
25%	3.210000	0.550000	9.500000	5.000000
50%	3.310000	0.620000	10.200000	6.000000
75%	3.400000	0.730000	11.100000	6.000000
max	4.010000	2.000000	14.900000	8.000000

FINDING NULL VALUES

```
df.isnull().any()
```

```
fixed acidity      False
volatile acidity   False
citric acid        False
residual sugar     False
chlorides          False
free sulfur dioxide False
total sulfur dioxide False
density            False
pH                 False
sulphates          False
alcohol            False
```

```
quality          False
dtype: bool
```

No Null values(so there are no missing to be handled)

```
df['quality'].value_counts()
5      611
6      575
7      186
4       46
8       17
3        6
Name: quality, dtype: int64
```

Corelation

```
df.corr()
```

	fixed acidity	volatile acidity	citric acid \
fixed acidity	1.000000	-0.256131	0.671703
volatile acidity	-0.256131	1.000000	-0.552496
citric acid	0.671703	-0.552496	1.000000
residual sugar	0.114777	0.001918	0.143577
chlorides	0.093705	0.061298	0.203823
free sulfur dioxide	-0.153794	-0.010504	-0.060978
total sulfur dioxide	-0.113181	0.076470	0.035533
density	0.668047	0.022026	0.364947
pH	-0.682978	0.234937	-0.541904
sulphates	0.183006	-0.260987	0.312770
alcohol	-0.061668	-0.202288	0.109903
quality	0.124052	-0.390558	0.226373

	residual sugar	chlorides	free sulfur
dioxide \			
fixed acidity	0.114777	0.093705	-0.153794
volatile acidity	0.001918	0.061298	-0.010504
citric acid	0.143577	0.203823	-0.060978
residual sugar	1.000000	0.055610	0.187049
chlorides	0.055610	1.000000	0.005562
free sulfur dioxide	0.187049	0.005562	1.000000
total sulfur dioxide	0.203028	0.047400	0.667666

density	0.355283	0.200632	-0.021946
pH	-0.085652	-0.265026	0.070377
sulphates	0.005527	0.371260	0.051658
alcohol	0.042075	-0.221141	-0.069408
quality	0.013732	-0.128907	-0.050656

	total sulfur dioxide	density	pH
sulphates \			
fixed acidity	-0.113181	0.668047	-0.682978
0.183006			
volatile acidity	0.076470	0.022026	0.234937
0.260987			-
citric acid	0.035533	0.364947	-0.541904
0.312770			
residual sugar	0.203028	0.355283	-0.085652
0.005527			
chlorides	0.047400	0.200632	-0.265026
0.371260			
free sulfur dioxide	0.667666	-0.021946	0.070377
0.051658			
total sulfur dioxide	1.000000	0.071269	-0.066495
0.042947			
density	0.071269	1.000000	-0.341699
0.148506			
pH	-0.066495	-0.341699	1.000000
0.196648			-
sulphates	0.042947	0.148506	-0.196648
1.000000			
alcohol	-0.205654	-0.496180	0.205633
0.093595			
quality	-0.185100	-0.174919	-0.057731
0.251397			

	alcohol	quality
fixed acidity	-0.061668	0.124052
volatile acidity	-0.202288	-0.390558
citric acid	0.109903	0.226373
residual sugar	0.042075	0.013732
chlorides	-0.221141	-0.128907
free sulfur dioxide	-0.069408	-0.050656
total sulfur dioxide	-0.205654	-0.185100
density	-0.496180	-0.174919
pH	0.205633	-0.057731
sulphates	0.093595	0.251397



```
alcohol      1.000000  0.476166
quality      0.476166  1.000000
```

```
df.corr()['quality'].sort_values(ascending=False)
```

```
quality      1.000000
alcohol      0.482143
sulphates    0.343689
citric acid   0.262596
fixed acidity 0.135510
residual sugar 0.010447
free sulfur dioxide -0.026502
pH            -0.084257
density       -0.175625
total sulfur dioxide -0.178415
chlorides     -0.194465
volatile acidity -0.395515
Name: quality, dtype: float64
```

```
sns.heatmap(df.corr(),annot=True)
```

```
<Axes: >
```




```
df['quality'] = df.quality.apply(lambda x : 1 if x > 6.5 else 0)
```

OUTLIER DETECTION


```
sns.boxplot(df["fixed acidity"])
```

<Axes: >




```
sns.boxplot(df["volatile acidity"])
```

<Axes: >




```
sns.boxplot(df["citric acid"])
```

<Axes: >




```
sns.boxplot(df["residual sugar"])
```

<Axes: >




```
sns.boxplot(df["chlorides"])
```

<Axes: >




```
sns.boxplot(df['free sulfur dioxide'])
```

<Axes: >




```
sns.boxplot(df["total sulfur dioxide"])
```

<Axes: >




```
sns.boxplot(df["density"])
```

<Axes: >




```
sns.boxplot(df["pH"])
```

```
<Axes: >
```




```
sns.boxplot(df["sulphates"])
```

```
<Axes: >
```



```
sns.boxplot(df["alcohol"])
```

<Axes: >



#Replacing the Outliers by IQR

```
# removal of outliers of "fixed acidity" by IQR
q1 = df['fixed acidity'].quantile(0.25)
```

```

q3 = df['fixed acidity'].quantile(0.75)
IQR = q3-q1
IQR

1.8000000000000007

upper_limit = q3+1.5*IQR
upper_limit

11.600000000000001


lower_limit =q1-1.5*IQR
lower_limit

4.399999999999999

df = df[df['fixed acidity']<upper_limit]
sns.boxplot(df["fixed acidity"] )

<Axes: >

```



```

# removal of outliers of "volatile acidity" by IQR
q1 = df['volatile acidity'].quantile(0.25)
q3 = df['volatile acidity'].quantile(0.75)
IQR = q3-q1
IQR

0.24375000000000002

```



```

upper_limit = q3+1.5*IQR
upper_limit

0.9993750000000001


lower_limit =q1-1.5*IQR
lower_limit

0.02437499999999998

df = df[df['volatile acidity']<upper_limit]
sns.boxplot(df["volatile acidity"])

<Axes: >

```



```

# removal of outliers of "citric acid" by IQR
q1 = df['citric acid'].quantile(0.25)
q3 = df['citric acid'].quantile(0.75)
IQR = q3-q1
IQR

0.31999999999999995

upper_limit = q3+1.5*IQR
upper_limit

0.8899999999999999


lower_limit =q1-1.5*IQR
lower_limit

```

```
-0.3899999999999999
```

```
df = df[df['citric acid']<upper_limit]  
sns.boxplot(df["citric acid"])
```

```
<Axes: >
```



```
# removal of outliers of "residual sugar" by IQR
```

```
q1 = df['residual sugar'].quantile(0.25)
```

```
q3 = df['residual sugar'].quantile(0.75)
```

```
IQR = q3-q1
```

```
IQR
```

```
0.5
```

```
upper_limit = q3+1.5*IQR
```

```
upper_limit
```

```
3.15
```

```
lower_limit =q1-1.5*IQR
```


```
lower_limit
```

```
1.15
```

```
df = df[df['residual sugar']<upper_limit]
```

```
sns.boxplot(df["residual sugar"])
```

```
<Axes: >
```



```
# removal of outliers of "chlorides" by IQR
```

```
q1 = df['chlorides'].quantile(0.25)
```

```
q3 = df['chlorides'].quantile(0.75)
```

```
IQR = q3-q1
```

```
IQR
```

```
0.016999999999999987
```

```
upper_limit = q3+1.5*IQR
```

```
upper_limit
```

```
0.11149999999999997
```

```
lower_limit =q1-1.5*IQR
```


```
lower_limit
```

```
0.0435000000000000025
```

```
df = df[df['chlorides']<upper_limit]
```

```
sns.boxplot(df["chlorides"])
```

```
<Axes: >
```



```
# removal of outliers of "free sulfur dioxide" by IQR
```

```
q1 = df['free sulfur dioxide'].quantile(0.25)
```

```
q3 = df['free sulfur dioxide'].quantile(0.75)
```

```
IQR = q3-q1
```

```
IQR
```

```
13.0
```

```
upper_limit = q3+1.5*IQR
```

```
upper_limit
```

```
40.5
```

```
lower_limit =q1-1.5*IQR
```


```
lower_limit
```

```
-11.5
```

```
df = df[df['free sulfur dioxide']<upper_limit]
```

```
sns.boxplot(df["free sulfur dioxide"])
```

```
<Axes: >
```



```
# removal of outliers of "total sulfur dioxide" by IQR
q1 = df['total sulfur dioxide'].quantile(0.25)
q3 = df['total sulfur dioxide'].quantile(0.75)
IQR = q3-q1
IQR

32.0

upper_limit = q3+1.5*IQR
upper_limit


101.0

lower_limit =q1-1.5*IQR
lower_limit

-27.0

df = df[df['total sulfur dioxide']<upper_limit]
sns.boxplot(df["total sulfur dioxide"])

<Axes: >
```



```
# removal of outliers of "density" by IQR
```

```
q1 = df['density'].quantile(0.25)
```

```
q3 = df['density'].quantile(0.75)
```

```
IQR = q3-q1
```

```
IQR
```

```
0.0021999999999999797
```

```
upper_limit = q3+1.5*IQR
```

```
upper_limit
```

```
1.0011
```

```
lower_limit =q1-1.5*IQR
```


```
lower_limit
```

```
0.992300000000000001
```

```
df = df[df['density']<upper_limit]
```

```
sns.boxplot(df.density)
```

```
<Axes: >
```



```
# removal of outliers of "pH" by IQR
```

```
q1 = df['pH'].quantile(0.25)
```

```
q3 = df['pH'].quantile(0.75)
```

```
IQR = q3-q1
```

```
IQR
```

```
0.18999999999999995
```

```
upper_limit = q3+1.5*IQR
```

```
upper_limit
```

```
3.6849999999999996
```

```
lower_limit =q1-1.5*IQR
```


```
lower_limit
```

```
2.925
```

```
df = df[df['pH']<upper_limit]
```

```
sns.boxplot(df["pH"])
```

```
<Axes: >
```



```
# removal of outliers of "sulphates" by IQR
```

```
q1 = df['sulphates'].quantile(0.25)
```

```
q3 = df['sulphates'].quantile(0.75)
```

```
IQR = q3-q1
```

```
IQR
```

```
0.15249999999999986
```

```
upper_limit = q3+1.5*IQR
```

```
upper_limit
```

```
0.9287499999999997
```

```
lower_limit =q1-1.5*IQR
```


```
lower_limit
```

```
0.321250000000000026
```

```
df = df[df['sulphates']<upper_limit]
```

```
sns.boxplot(df["sulphates"])
```

```
<Axes: >
```

```
# removal of outliers of "sulphates" by IQR
```

```
q1 = df['alcohol'].quantile(0.25)
```

```
q3 = df['alcohol'].quantile(0.75)
```

```
IQR = q3-q1
```

```
IQR
```

```
1.5
```

```
upper_limit = q3+1.5*IQR
```

```
upper_limit
```

```
13.25
```

```
lower_limit =q1-1.5*IQR
```


```
lower_limit
```

```
7.25
```

```
df = df[df['alcohol']<upper_limit]
```

```
sns.boxplot(df["alcohol"])
```

```
<Axes: >
```



SPLITTING THE VALUES TO X AND Y

```
y=df['quality']
y
0      0
1      0
2      0
3      0
4      0
..
1594    0
1595    0
1596    0
1597    0
1598    0
Name: quality, Length: 1599, dtype: int64
```

```
x = df.drop(columns=['quality'],axis=1)
x.head()
```

	fixed acidity	volatile acidity	citric acid	residual sugar
0	7.4	0.70	0.00	1.9
1	7.8	0.88	0.00	2.6
2	7.8	0.76	0.04	2.3

```

0.092
3      11.2      0.28      0.56      1.9
0.075
4      7.4      0.70      0.00      1.9
0.076

    free sulfur dioxide  total sulfur dioxide  density  pH  sulphates
\
0      11.0      34.0  0.9978  3.51      0.56
1      25.0      67.0  0.9968  3.20      0.68
2      15.0      54.0  0.9970  3.26      0.65
3      17.0      60.0  0.9980  3.16      0.58
4      11.0      34.0  0.9978  3.51      0.56

    alcohol
0      9.4
1      9.8
2      9.8
3      9.8
4      9.4

```

SCALING

```

from sklearn.preprocessing import MinMaxScaler
scale =MinMaxScaler()

x_scaled= pd.DataFrame(scale.fit_transform(x),columns =x.columns)
x_scaled.head()

    fixed acidity  volatile acidity  citric acid  residual sugar
chlorides \
0      0.247788      0.397260      0.00      0.068493
0.106845
1      0.283186      0.520548      0.00      0.116438
0.143573
2      0.283186      0.438356      0.04      0.095890
0.133556
3      0.584071      0.109589      0.56      0.068493
0.105175
4      0.247788      0.397260      0.00      0.068493
0.106845

    free sulfur dioxide  total sulfur dioxide  density  pH
sulphates \

```

```

0      0.140845      0.098940  0.567548  0.606299
0.137725
1      0.338028      0.215548  0.494126  0.362205
0.209581
2      0.197183      0.169611  0.508811  0.409449
0.191617
3      0.225352      0.190813  0.582232  0.330709
0.149701
4      0.140845      0.098940  0.567548  0.606299
0.137725

      alcohol
0  0.153846
1  0.215385
2  0.215385
3  0.215385
4  0.153846

```

TRAIN TEST SPLIT

```

from sklearn.model_selection import train_test_split
x_train,x_test,y_train,y_test =
train_test_split(x_scaled,y,test_size=0.3,random_state=10)

x_train.shape
(1119, 11)

x_train.head()

```

	fixed acidity	volatile acidity	citric acid	residual sugar
chlorides \				
305	0.504425	0.280822	0.48	0.109589
0.085142				
984	0.672566	0.226027	0.49	0.034247
0.105175				
47	0.362832	0.116438	0.52	0.047945
0.168614				
812	0.548673	0.226027	0.33	0.109589
0.145242				
3	0.584071	0.109589	0.56	0.068493
0.105175				

	free sulfur dioxide	total sulfur dioxide	density	pH
free sulphates \				
305	0.070423	0.067138	0.714391	0.299213
0.155689				
984	0.028169	0.000000	0.501468	0.307087
0.179641				

47	0.154930	0.109541	0.501468	0.401575
0.149701				
812	0.267606	0.113074	0.595448	0.393701
0.227545				
3	0.225352	0.190813	0.582232	0.330709
0.149701				

	alcohol
305	0.138462
984	0.307692
47	0.169231
812	0.369231
3	0.215385

y_train.shape

(1119,)

y_train.head()

305	0
984	0
47	0
812	0
3	0

Name: quality, dtype: int64

x_test.shape

(480, 11)

x_test.head()

	fixed acidity	volatile acidity	citric acid	residual sugar
chlorides \				
1518	0.247788	0.239726	0.46	0.089041
0.170284				
1246	0.247788	0.424658	0.07	0.054795
0.123539				
544	0.858407	0.130137	0.74	0.061644
0.105175				
1343	0.256637	0.267123	0.02	0.054795
0.120200				
428	0.398230	0.273973	0.33	0.027397
0.096828				

	free sulfur dioxide	total sulfur dioxide	density	pH	\
1518	0.084507	0.049470	0.469897	0.456693	
1246	0.197183	0.148410	0.363436	0.299213	
544	0.070423	0.031802	0.787812	0.094488	
1343	0.169014	0.088339	0.389868	0.488189	

```
428          0.112676          0.084806  0.567548  0.393701
```

```
      sulphates  alcohol
1518    0.179641  0.323077
1246    0.089820  0.246154
544     0.275449  0.000000
1343    0.125749  0.323077
428     0.161677  0.138462
```

```
y_test.shape
```

```
(480,)
```

```
y_test.head()
```

```
1518    0
1246    0
544     0
1343    0
428     0
Name: quality, dtype: int64
```

MODEL BUILDING

```
from sklearn.linear_model import LogisticRegression
```

```
model = LogisticRegression()
model.fit(x_train,y_train)
```

```
LogisticRegression()
```

```
y_pred =model.predict(x_test)
y_pred
```

```
array([0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
0,
      0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0,
0,
      0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
0,
      0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
0,
      0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
0,
      0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
0,
      0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
0,
      0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 1, 1, 0, 0, 0, 1, 0, 0, 0,
```

```

0,
    0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0,
0,
    0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0,
0,
    0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
0,
    0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0,
0,
    0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
0,
    0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
1,
    0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
0,
    0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
0,
    1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
0,
    0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
0,
    0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
0,
    0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
0,
    0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
0,
    0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0])

```

```

y_pred = model.predict(x_train)
y_pred

```

```

array([0, 0, 0, ..., 0, 0, 0])

```

```

from sklearn.metrics import
accuracy_score, classification_report, confusion_matrix

```

```

import sklearn.metrics as metrics

```

```

y_pred = model.predict(x_test)
print("Classification Report", metrics.classification_report(y_test,
y_pred))
print("Accuracy score:", metrics.accuracy_score(y_test, y_pred))
print("Precision score:", metrics.precision_score(y_test, y_pred,
average = 'macro'))
print("Recall score:", metrics.recall_score(y_test, y_pred, average =
'macro'))

```

```

Classification Report          precision    recall  f1-score
support

```

	0	0.88	0.99	0.93	411
	1	0.72	0.19	0.30	69
accuracy				0.87	480
macro avg		0.80	0.59	0.61	480
weighted avg		0.86	0.87	0.84	480

Accuracy score: 0.8729166666666667
Precision score: 0.8005050505050505
Recall score: 0.5881201734898974

```
confusion_matrix(y_test,y_pred)
```

```
array([[406,  5],
       [ 56, 13]])
```

```
pd.crosstab(y_test,y_pred)
```

```
col_0    0    1
quality
0        406    5
1         56   13
```

```
print(classification_report(y_test,y_pred))
```

	precision	recall	f1-score	support
0	0.88	0.99	0.93	411
1	0.72	0.19	0.30	69
accuracy			0.87	480
macro avg	0.80	0.59	0.61	480
weighted avg	0.86	0.87	0.84	480

MAX ACCURACY SCORE IS 87

DECISION TREE CLASSIFIER

```
from sklearn.tree import DecisionTreeClassifier

model1 =
DecisionTreeClassifier(max_depth=4,splitter='best',criterion='entropy'
)

model1.fit(x_train,y_train)

DecisionTreeClassifier(criterion='entropy', max_depth=4)
```



```

d_y_predict = model1.predict(x_test)
d_y_predict
array([0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
0,
      0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0,
0,
      0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0,
0,
      0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0,
0,
      0, 0, 1, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
0,
      0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
0,
      0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1,
1,
      0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 1, 0, 0, 0, 1, 0, 0, 0,
0,
      0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0,
0,
      0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0,
0,
      0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
0,
      0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
0,
      0, 1, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 1,
0,
      0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0,
0,
      0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
0,
      0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0,
0,
      0, 1, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 1, 0,
0,
      1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
0,
      0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 1, 0,
0,
      0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0,
0,
      1, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
0,
      0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 1, 0, 0,
0,
      0, 1, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0])

d_y_predict_train = model1.predict(x_train)
d_y_predict_train

```

```
array([0, 0, 0, ..., 0, 0, 0])

from sklearn.metrics import
accuracy_score,classification_report,confusion_matrix

print('Testing Accuracy = ', accuracy_score(y_test,d_y_predict))
print('Training Accuracy = ',
accuracy_score(y_train,d_y_predict_train))
```

```
Testing Accuracy = 0.8541666666666666
Training Accuracy = 0.8972296693476318
```

```
pd.crosstab(y_test,d_y_predict)
```

col_0	0	1
quality		
0	390	21
1	49	20

```
print(classification_report(y_test,d_y_predict))
```

	precision	recall	f1-score	support
0	0.89	0.95	0.92	411
1	0.49	0.29	0.36	69
accuracy			0.85	480
macro avg	0.69	0.62	0.64	480
weighted avg	0.83	0.85	0.84	480

MAX ACCURACY SCORE IS 85

RANDOM FOREST CLASSIFIER

```
from sklearn.ensemble import RandomForestClassifier
model2 = RandomForestClassifier(n_estimators=200, criterion='entropy')
model2.fit(x_train, y_train)

RandomForestClassifier(criterion='entropy', n_estimators=200)

r_y_predict = model2.predict(x_test)
r_y_predict_train = model2.predict(x_train)
r_y_predict

array([0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0,
0,
```

```

0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 1, 0, 0, 0, 0, 0, 0,
0, 0, 0, 0, 1, 0, 0, 0, 1, 0, 1, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0,
0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 1, 0, 0, 0, 0, 0, 0,
0, 0, 1, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0,
0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0,
0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0,
0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0,
1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0,
0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 1, 1, 0, 0, 0, 0, 0,
0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0,
0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0,
0, 0, 0, 0, 0, 0, 0, 1, 0, 1, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0,
0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
0, 1, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0,
0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0,
0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 1, 0, 0,
0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
0, 1, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0])

```

```

r_y_predict_train

```

```

array([0, 0, 0, ..., 0, 0, 0])

```

```

print('Testing Accuracy = ', accuracy_score(y_test,r_y_predict))

```

```

print('Training Accuracy = ',
accuracy_score(y_train,r_y_predict_train))

```

```
Testing Accuracy = 0.9083333333333333
```

```
Training Accuracy = 1.0
```

```
pd.crosstab(y_test,r_y_predict)
```

```
col_0      0      1
quality
0         404      7
1          37     32
```

```
print(classification_report(y_test,r_y_predict))
```

	precision	recall	f1-score	support
0	0.92	0.98	0.95	411
1	0.82	0.46	0.59	69
accuracy			0.91	480
macro avg	0.87	0.72	0.77	480
weighted avg	0.90	0.91	0.90	480

MAX ACCURACY SCORE IS 91

GOT THE HIGHEST ACCURACY SCORE IN RANDOM FOREST CLASSIFIER

Test with random observation

```
x.head()
```

	fixed acidity	volatile acidity	citric acid	residual sugar
0	7.4	0.70	0.00	1.9
0.076				
1	7.8	0.88	0.00	2.6
0.098				
2	7.8	0.76	0.04	2.3
0.092				
3	11.2	0.28	0.56	1.9
0.075				
4	7.4	0.70	0.00	1.9
0.076				

	free sulfur dioxide	total sulfur dioxide	density	pH	sulphates
0	11.0	34.0	0.9978	3.51	0.56
1	25.0	67.0	0.9968	3.20	0.68
2	15.0	54.0	0.9970	3.26	0.65
3	17.0	60.0	0.9980	3.16	0.58
4	11.0	34.0	0.9978	3.51	0.56

	alcohol
0	9.4
1	9.8
2	9.8
3	9.8
4	9.4


```
print("Prediction:",model.predict([[0.283186,0.520548,0.56,0.068493,0.143573,0.098940,0.567548,0.409449,0.137725,0.191617,0.215385]]))
```

Prediction: [0]

```
/usr/local/lib/python3.10/dist-packages/sklearn/base.py:439:
UserWarning: X does not have valid feature names, but
LogisticRegression was fitted with feature names
warnings.warn(
```



```
print("Prediction:",model.predict([[0.380531,0.109589,0.45,0.054795,0.091820,0.084507,0.021201,0.254772,0.401575,0.131737,0.600000]]))
```

Prediction: [0]

```
/usr/local/lib/python3.10/dist-packages/sklearn/base.py:439:
UserWarning: X does not have valid feature names, but
LogisticRegression was fitted with feature names
warnings.warn(
```

CONCLUSION : all the three models gave the "alcohol quality is BAD"